

1. K-Nearest Neighbors (KNN):

- How does the KNN algorithm classify a new data point?
- What is the significance of the "K" value in KNN?
- How does KNN handle categorical features?
- Can KNN be computationally expensive for large datasets?

2. Naive Bayes:

- What makes Naive Bayes "naive" in its assumptions?
- How does Naive Bayes handle missing data?
- Explain the concept of prior and posterior probabilities in the context of Naive Bayes.
- When is Naive Bayes particularly suitable for classification tasks?

3. Gradient Descent:

- What is the main goal of the gradient descent optimization algorithm?
- Differentiate between batch gradient descent and stochastic gradient descent.
- How does learning rate impact the convergence of gradient descent?
- Discuss the concept of local minima in the context of gradient descent.

4. Linear Regression:

- What is the fundamental equation of a linear regression model?
- Explain the terms "slope" and "intercept" in the context of linear regression.
- How is the cost function used in linear regression optimization?
- What assumptions does linear regression make about the data?

5. Neural Network:

- What distinguishes a neural network from traditional machine learning models?
- Define the terms "input layer," "hidden layer," and "output layer" in a neural network.
- How does backpropagation work in training a neural network?
- Discuss the concept of overfitting in neural networks.

6. Convolutional Neural Network (CNN):

- What is the primary advantage of using convolutional layers in CNNs for image processing?
- Explain the purpose of pooling layers in a CNN.
- How do CNNs handle spatial hierarchies in image data?
- Discuss the concept of weight sharing in convolutional layers.